2026

4 Nov — Literature review: 29 papers, MUCS leaderboard, implementation brief

7 Nov — Stage 1A: data pipeline, tokenizer, smoke test, GPU profiling

8 Nov — Iteration 1: baseline fine-tuning (38,000 updates, 8 h 52 min)

8 Nov — Iteration 2: continued training from iteration 1 (6 h 59 min)

9 Nov — Iteration 3: script-normalised targets (6 h 59 min)

9 Nov — Iteration 4: misalignment filter (7 h 02 min)

selected model

9 Nov — Iteration 5: CTC realignment of training segments (6 h 52 min)

10 Nov — CPU inference backend, ARPA LM reader, Streamlit dashboard

after the experiments